



Airflow Cheat Sheet

Airflow CLI (Command Line Interface)

GENERAL COMMANDS

- **airflow version** :Shows the version of Airflow
- **airflow info** :Provides sys. and dependencies info.
- **airflow webserver** :Starts Airflow webserver
- **airflow scheduler** :Starts the Airflow scheduler
- **airflow worker** :Starts a new Airflow worker

DAG COMMANDS

- **airflow dags list** :List all the DAGs
- **airflow dags delete DAG_ID** :Deletes a DAG
- **airflow dags trigger -d DAG_ID** :Manually triggers a DAG run
- **airflow dags test DAG_ID EXECUTION_DATE** :Runs a specific DAG
- **airflow dags pause DAG_ID** :Pauses a DAG
- **airflow dags unpause DAG_ID** :Unpauses a DAG

TASK COMMANDS

- **airflow tasks list DAG_ID** :List all the tasks in a DAG
- **airflow tasks test DAG_ID TASK_ID EXECUTION_DATE** :Test a task instance. This will not produce any database entries or trigger handlers
- **airflow tasks run DAG_ID TASK_ID EXECUTION_DATE** :Run a single task instance
- **airflow tasks clear -t TASK_REGEX -s START_DATE -e END_DATE DAG_ID** :Clears the task instances matching the filters from the history.


POOL COMMANDS

- **airflow pools list** :List all pools
- **airflow pools get POOL_NAME** :Get pool size and used slots count
- **airflow pools set POOL_NAME POOL_SLOT_COUNT POOL_DESCRIPTION** :Create a new pool with the given parameters.
- **airflow pools delete POOL_NAME** :Delete a pool

CONNECTION COMMANDS

- **airflow connections list** :List all connections
- **airflow connections add 'CONN_ID' --conn-type 'CONN_TYPE' --conn-host 'CONN_HOST' --conn-login 'CONN_LOGIN' --conn-password 'CONN_PASSWORD' --conn-schema 'CONN_SCHEMA' --conn-port 'CONN_PORT'** :Add new connection
- **airflow connections delete 'CONN_ID'** :Delete a connection





VARIABLE COMMANDS

- `airflow variables list` :List all variables
- `airflow variables get 'VAR_KEY'` :Get value of a variable
- `airflow variables set 'VAR_KEY' 'VAR_VALUE'` :Set a variable and its value
- `airflow variables delete 'VAR_KEY'` :Delete a variable

Creating a DAG (Directed Acyclic Graph)

BASIC DAG DECLARATION

```


• from airflow import DAG
  from airflow.operators.dummy_operator import DummyOperator
  from datetime import datetime, timedelta
  default_args = {
      'owner': 'airflow',
      'depends_on_past': False,
      'start_date': datetime(2022, 1, 1),
      'email': ['airflow@example.com'],
      'email_on_failure': False,
      'email_on_retry': False,
      'retries': 1,
      'retry_delay': timedelta(minutes=5),
  }
  with DAG('my_first_dag', default_args=default_args, description='A simple tutorial DAG',
    schedule_interval=timedelta(days=1),) as dag:
      t1 = DummyOperator(task_id='task_1')
      t2 = DummyOperator(task_id='task_2')
      t1 >> t2

```

DAG ARGUMENTS

- | | |
|--|--|
| <ul style="list-style-type: none"> • owner :Owner of the DAG • start_date :The start date of the DAG • email :The email(s) to notify on task failure • depends_on_past :If True, a task can't run unless the previous schedule succeeded • email_on_failure :If True, emails are sent on task failure • email_on_retry :If True, emails are sent on each task retry • description :Description of the DAG • schedule_interval :The interval of DAG execution | <ul style="list-style-type: none"> • retries :Number of times a failed task is retried • retries :Number of times a failed task is retried • retry_delay :Time delay between retries |
|--|--|





ORDERING AND DEPENDENCIES

- `t1 >> t2` :t1 precedes t2
- `t1 << t2` :Get value of a variable
- `chain(t1, t2, t3)` :Chains the tasks in sequence (t1 >> t2 >> t3).
- `cross_downstream([t1, t2], [t3])` :Creates dependencies t1 >> t3 and t2 >> t3
- `t1.set_upstream(t2)` : t1 follows t2.
- `t1.set_downstream(t2)` :t1 precedes t2

Using Operators

BASIC OPERATOR USAGE

- ```
from airflow import DAG
from airflow.operators.dummy_operator import DummyOperator
from datetime import datetime
with DAG('my_dag', start_date=datetime(2022, 1, 1)) as dag:
 task = DummyOperator(task_id='my_task')
```

### COMMON OPERATORS AND THEIR PARAMETERS

- **DummyOperator** :Does nothing, but useful for testing or if you need a no-op operator. Its main parameter is `task_id`  
`DummyOperator(task_id='dummy_task')`
- **PythonOperator** :Executes a Python function. Main parameters are `task_id` and `python_callable` (the function to run)  

```
from airflow.operators.python_operator import PythonOperator
def my_func():
 print('Hello, World!')
PythonOperator(task_id='python_task', python_callable=my_func)
```
- **BashOperator** :Executes a bash command. Main parameters are `task_id` and `bash_command`  
`BashOperator(task_id='bash_task', bash_command='echo Hello, World!')`
- **EmailOperator** :Sends an email. Main parameters are `task_id`, `to` (recipient address), `subject`, and `html_content`  

```
from airflow.operators.email_operator import EmailOperator
EmailOperator(task_id='email_task', to='test@example.com', subject='Hello',
html_content='<h1>World!</h1>')
```
- **SubDagOperator** :Executes a sequence of tasks in a sub-DAG.  

```
from airflow.operators.subdag_operator import SubDagOperator
SubDagOperator(task_id='subdag_task', subdag=subdag())
```

### CREATING A CUSTOM OPERATOR

Creating a custom operator involves defining a new Python class that inherits from the BaseOperator class.

- ```
from airflow.models import BaseOperator
from airflow.utils.decorators import apply_defaults
class MyFirstOperator(BaseOperator):
    @apply_defaults
```





```
def __init__(self, my_param, *args, **kwargs):
    super(MyFirstOperator, self).__init__(*args, **kwargs)
    self.my_param = my_param
def execute(self, context):
    log.info('My first operator, param: %s', self.my_param)
```

In this example, `MyFirstOperator` takes a parameter `my_param`. The `execute` method is called when the operator is run.

Tasks

BASIC TASK DECLARATION

- A task is an instance of an operator. Here's a basic task declaration using the `PythonOperator`:

```
from airflow.operators.python_operator import PythonOperator
def hello_world():
    print('Hello, World!')
task = PythonOperator(task_id='hello_world_task', python_callable=hello_world)
```

In this example, `task_id` is a unique identifier for the task, and `python_callable` is the Python function that the task will execute


TASK ARGUMENTS

- `depends_on_past`** : If set to True, the task instance will only run if the previous task instance succeeded. Defaults to False
- `wait_for_downstream`** : If set to True, this task instance will wait for tasks downstream of the previous task instance to finish before running.
- `retries`** : The number of times to retry failed tasks. Defaults to 0.
- `retry_delay`** : The time to wait before retrying the task. Defaults to five minutes. This must be a `datetime.timedelta` object.
- `queue`** : The name of the queue to use for task execution. This is used with the `CeleryExecutor` and `KubernetesExecutor`.
- `priority_weight`** : Defines the priority of the task in relation to other tasks. This can be used to prioritize certain tasks in the scheduler.
- `pool`** : The name of the pool to use, if any. This can be used to limit parallelism for a set of tasks.
- `trigger_rule`** : Defines the rule by which the task gets triggered (`all_success`, `all_failed`, `all_done`, `one_success`, `one_failed`, `none_failed`, `none_skipped`, etc.)

Example:-

- `task = PythonOperator(`
`task_id='hello_world_task',`
`python_callable=hello_world,`
`depends_on_past=True,`
`wait_for_downstream=True,`
`retries=3,`
`retry_delay=timedelta(minutes=1),`
`queue='my_queue',`
`priority_weight=10,`
`pool='my_pool'`
`)`





TaskFlow API

THE @task DECORATOR

- `from airflow.decorators import task`
`@task`
`def my_task():`
 `print('Hello, World!')`

DYNAMIC OUTPUT TYPES

- `from airflow.decorators import task`
`from typing import Tuple`
`@task`
`def my_task() -> Tuple[int, str]:`
 `return 42, 'Hello, World!'`

USING XComArg FOR INTER-TASK COMMUNICATION

- `from airflow.decorators import task`
`from airflow.models import XComArg`
`@task`
`def my_task() -> str:`
 `return 'Hello, World!'`
`@task`
`def greet(msg: XComArg):`
 `print(f"Greeting: {msg}")`
`#Use my_task's output in greet task`
`greet(my_task())`

Conditional Execution and Branching


The BranchPythonOperator

- The BranchPythonOperator is used when you want to follow different paths in your DAG based on a condition. It takes a Python callable that returns a task id (or a list of task ids).
- `from airflow.operators.python import BranchPythonOperator`
`def choose_path():`
 `return 'first_task' if CONDITION else 'second_task'`
`branch_task = BranchPythonOperator(`
 `task_id='branch_task',`
 `python_callable=choose_path`
`)`
`#Define first_task and second_task here`

The ShortCircuitOperator

- The ShortCircuitOperator is used to bypass (short-circuit) a section of the DAG based on a condition. It is a PythonOperator that evaluates a condition and short-circuits the workflow if the condition is False.
- `from airflow.operators.python import ShortCircuitOperator`
`def check_condition():`
 `return True if CONDITION else False`
`condition_task = ShortCircuitOperator(`
 `task_id='condition_task',`
 `python_callable=check_condition`
`)`
`#Define further tasks here. They won't run if check_condition returns False`





Dynamic DAGs

CREATING DYNAMIC DAGs

```
• from airflow import DAG
  from airflow.operators.dummy_operator import
    DummyOperator
  from airflow.utils.dates import days_ago
  #Function to create a DAG
  def create_dag(dag_id, schedule):
      dag = DAG(
          dag_id=dag_id,
          schedule_interval=schedule,
          start_date=days_ago(2),
          default_args={'retries': 1},
      )
      with dag:
          DummyOperator(task_id='task_1')
          DummyOperator(task_id='task_2')
      return dag
  #Loop to create multiple DAGs
  for i in range(1, 6):
      dag_id = f"my_dag_{i}"
      globals()[dag_id] = create_dag(dag_id,
                                      '@daily')
```

DYNAMIC TASK GENERATION

```
• from airflow import DAG
  from airflow.operators.dummy_operator import
    DummyOperator
  from airflow.utils.dates import days_ago
  dag = DAG(
      'my_dag',
      start_date=days_ago(2),
      default_args={'retries': 1},
  )
  #Loop to create multiple tasks
  for i in range(1, 6):
      DummyOperator(task_id=f'task_{i}', dag=dag)
```

Error Handling


CALLBACKS

- **on_failure_callback** : Function called on task failure. Use for custom failure behavior (e.g., notifications or cleanup).
- **on_retry_callback** : Function called before task retry. Use for pre-retry adjustments or logging.
- **on_success_callback** : Function called on task success. Use for post-success actions (e.g., notifications or downstream triggers).

Example:-

```
#Define your callback functions
def task_fail_callback(context): #Write body to handle failure below
def task_retry_callback(context): #Write body to handle retry below
def task_success_callback(context): #Write body to handle success below
#Assign callbacks to your task
task = PythonOperator(
    task_id='task',
    python_callable=my_function,
    on_failure_callback=task_fail_callback,
    on_retry_callback=task_retry_callback,
    on_success_callback=task_success_callback,
    dag=dag,
)
```





TRIGGERRULE CLASS FOR HANDLING TASK DEPENDENCIES

- **TriggerRule.ALL_SUCCESS** : All parent tasks must have succeeded. This is the default rule.
- **TriggerRule.ALL_FAILED** : All parent tasks are in the 'failed' or 'upstream_failed' state.
- **TriggerRule.ALL_DONE** : All parent tasks are done with their execution.
- **TriggerRule.ONE_SUCCESS** : At least one parent task has succeeded.
- **TriggerRule.ONE_FAILED** : At least one parent task is in the 'failed' or 'upstream_failed' state.
- **TriggerRule.NONE_FAILED** : No parent task is in the 'failed' or 'upstream_failed' state.

Example:-

```
from airflow.utils.trigger_rule import TriggerRule
#Create tasks with different trigger rules
task1 = DummyOperator(task_id='task1', dag=dag)
task2 = DummyOperator(task_id='task2', dag=dag, trigger_rule=TriggerRule.ALL_SUCCESS)
task3 = DummyOperator(task_id='task3', dag=dag, trigger_rule=TriggerRule.ONE_FAILED)
task4 = DummyOperator(task_id='task4', dag=dag, trigger_rule=TriggerRule.ALL_DONE)
#Setting dependencies
task1 >> task2 >> [task3, task4]
```

Integration with Other Tools

EXAMPLE FOR FILE TRANSFER OPERATORS

• S3ToRedshiftOperator

```
from airflow.providers.amazon.aws.transfers.s3_to_redshift import S3ToRedshiftOperator
transfer_s3_redshift = S3ToRedshiftOperator(
    task_id='transfer_s3_redshift',
    schema='SCHEMA_NAME',
    table='TABLE_NAME',
    s3_bucket='S3_BUCKET',
    s3_key='S3_KEY',
    copy_options=['csv'],
    dag=dag,
)
```

• GoogleCloudStorageToBigQueryOperator

```
from airflow.providers.google.cloud.transfers.gcs_to_bigquery import GoogleCloudStorageToBigQueryOperator
transfer_gcs_bq = GoogleCloudStorageToBigQueryOperator(
    task_id='gcs_to_bq_example',
    bucket='BUCKET_NAME',
    source_objects=['SOURCE_FILE_NAME.csv'],
    destination_project_dataset_table='DESTINATION_DATASET.TABLE_NAME',
    schema_fields=[{'name': 'column1', 'type': 'STRING'}, {'name': 'column2', 'type': 'STRING'}],
    write_disposition='WRITE_TRUNCATE',
    dag=dag,
)
```





EXAMPLE FOR DATABASE OPERATORS

• PostgresOperator

```
from airflow.providers.postgres.operators.postgres import PostgresOperator
t1 = PostgresOperator(
    task_id='run_sql',
    sql='sql_file.sql',
    postgres_conn_id='postgres_default',
    dag=dag,
)
```

• MySQLOperator


```
from airflow.providers.mysql.operators.mysql import MySQLOperator
sql_op = MySQLOperator(
    task_id='run_sql',
    sql='sql_file.sql',
    mysql_conn_id='mysql_default',
    dag=dag,
)
```


Testing and Debugging

- **TESTING A TASK :** `airflow tasks test DAG_ID TASK_ID EXECUTION_DATE`
- **BACKFILLING DATA :** `airflow dags backfill -s START_DATE -e END_DATE DAG_ID`
- **RUNNING A TASK INSTANCE :** `airflow tasks run DAG_ID TASK_ID EXECUTION_DATE`

TriggerRule Class

- **all_success** : All parent tasks must have succeeded. This is the default trigger rule.
`task3 = DummyOperator(task_id='task3', trigger_rule=TriggerRule.ALL_SUCCESS, dag=dag)`
- **all_failed** : All parent tasks must have failed
`task3 = DummyOperator(task_id='task3', trigger_rule=TriggerRule.ALL_FAILED, dag=dag)`
- **all_done** : All parent tasks are done with their execution.
`task3 = DummyOperator(task_id='task3', trigger_rule=TriggerRule.ALL_DONE, dag=dag)`
- **one_success** : Fires if at least one parent has succeeded, it allows triggering the task soon after one parent succeeds.
`task3 = DummyOperator(task_id='task3', trigger_rule=TriggerRule.ONE_SUCCESS, dag=dag)`
- **one_failed** : Fires if at least one parent has failed, it allows triggering the task soon after one parent fails.
`task3 = DummyOperator(task_id='task3', trigger_rule=TriggerRule.ONE_FAILED, dag=dag)`
- **none_failed** : Fires as long as no parents have failed, i.e. all parents have succeeded or been skipped.
`task3 = DummyOperator(task_id='task3', trigger_rule=TriggerRule.NONE_FAILED, dag=dag)`





How to use **ChatGPT** Effectively for Airflow Workflow

ChatGPT and CodeGPT Prompts for Airflow

• SETTING UP A DAG IN APACHE AIRFLOW

- **ChatGPT:** "Explain the step-by-step process to setup a Directed Acyclic Graph (DAG) in Apache Airflow."
- **CodeGPT:** "Generate a simple DAG in Apache Airflow with 3 tasks that run in sequence."

• IMPLEMENTING CONDITIONAL EXECUTION IN APACHE AIRFLOW

- **ChatGPT:** "Detail the method for implementing conditional execution in Apache Airflow. What operators are commonly used?"
- **CodeGPT:** "Create an example DAG in Apache Airflow demonstrating conditional execution using the BranchPythonOperator."

• HANDLING ERRORS IN APACHE AIRFLOW

- **ChatGPT:** "How should I handle errors in Apache Airflow tasks? What are the best practices?"
- **CodeGPT:** "Generate a sample task in Apache Airflow with on_failure_callback, on_retry_callback, and on_success_callback."

• INTERACTING WITH DATABASES IN APACHE AIRFLOW

- **ChatGPT:** "Explain the process of setting up a connection to a PostgreSQL database in Apache Airflow."
- **CodeGPT:** "Generate a sample PostgresOperator task in an Airflow DAG that executes a simple SQL command."

• AIRFLOW WITH KUBERNETES

- **ChatGPT:** "Explain how to leverage Kubernetes in Airflow? What are the benefits and complexities involved?"
- **CodeGPT:** "Generate a simple KubernetesPodOperator task in an Airflow DAG."

• BACKFILLING DATA IN AIRFLOW

- **ChatGPT:** "How to backfill data in Airflow? What considerations should be kept in mind when doing so?"
- **CodeGPT:** "Show me an example of an Airflow command for backfilling a DAG from a specific date."

• SCALING AIRFLOW

- **ChatGPT:** "What strategies can be applied to scale Airflow? How do I ensure my Airflow setup remains robust as data volume increases?"
- **CodeGPT:** "Provide an example of how to set up the LocalExecutor in Airflow for better scalability."


• EXCEPTION HANDLING IN AIRFLOW

- **ChatGPT:** "Explain how to handle exceptions in Apache Airflow effectively. What are the best practices?"
- **CodeGPT:** "Provide a Python snippet demonstrating the use of on_failure_callback in Apache Airflow for exception handling."

• CUSTOM AIRFLOW PLUGINS

- **ChatGPT:** "How can I develop and use custom plugins in Airflow? What are the steps and important considerations?"
- **CodeGPT:** "Generate a skeleton Python code for a custom Airflow plugin."





• SECURE CONNECTIONS IN AIRFLOW

- **ChatGPT:** *"What are the best practices to handle secure connections in Airflow? How to use connection hooks?"*
- **CodeGPT:** *"Generate a Python snippet demonstrating the use of connection hooks for secure database connection."*

• AIRFLOW WITH AWS: S3 AND REDSHIFT

- **ChatGPT:** *"How can Apache Airflow be used with AWS services, such as S3 and Redshift? Describe with use-cases."*
- **CodeGPT:** *"Write an example DAG using **S3ToRedshiftOperator** to load data from S3 to Redshift."*

• AIRFLOW IN CI/CD PIPELINES

- **ChatGPT:** *"How can Apache Airflow be incorporated into a CI/CD pipeline? What are the benefits?"*
- **CodeGPT:** *"Generate a Python snippet showing an example DAG, which could be part of a CI/CD pipeline."*

• COMPLEX DEPENDENCY MANAGEMENT IN AIRFLOW

- **ChatGPT:** *"How to manage complex dependencies in Apache Airflow? Share tips for handling inter-DAG and intra-DAG dependencies."*
- **CodeGPT:** *"Generate a Python snippet showing an example DAG, which could be part of a CI/CD pipeline."*

• AIRFLOW WITH GOOGLE CLOUD PLATFORM SERVICES

- **ChatGPT:** *"How can Apache Airflow be integrated with Google Cloud Platform services, such as GCS and BigQuery? Describe use-cases."*
- **CodeGPT:** *"Generate a Python script showing an example DAG using **GoogleCloudStorageToBigQueryOperator** to load data from GCS to BigQuery."*

• ADVANCED ERROR HANDLING IN AIRFLOW

- **ChatGPT:** *"Explain strategies for advanced error handling in Apache Airflow, including fallbacks and failovers."*
- **CodeGPT:** *"Create a Python snippet demonstrating the use of **on_retry_callback** and **on_failure_callback** for advanced error handling."*

• AIRFLOW SCALABILITY AND HIGH AVAILABILITY

- **ChatGPT:** *"What strategies can be applied to ensure high availability and scalability of Apache Airflow deployments?"*
- **CodeGPT:** *"Create a sample Airflow configuration demonstrating optimal settings for a highly available and scalable deployment."*

• AIRFLOW WITH MACHINE LEARNING PIPELINES

- **ChatGPT:** *"How can Apache Airflow be leveraged for building and managing Machine Learning pipelines?"*
- **CodeGPT:** *"Write a Python script showing a simple example DAG for a machine learning pipeline using scikit-learn or TensorFlow."*

• MONITORING AND ALERTING IN AIRFLOW

- **ChatGPT:** *"What are the best practices for setting up monitoring and alerting for Apache Airflow workflows?"*
- **CodeGPT:** *"Create an example of how to set up a custom alert using **email_on_failure** and **email_on_retry** parameters in a task."*

